

Top 10 Kafka Use Cases



ASHISH PRATAP SINGH

MAR 27, 2025 · PAID

78

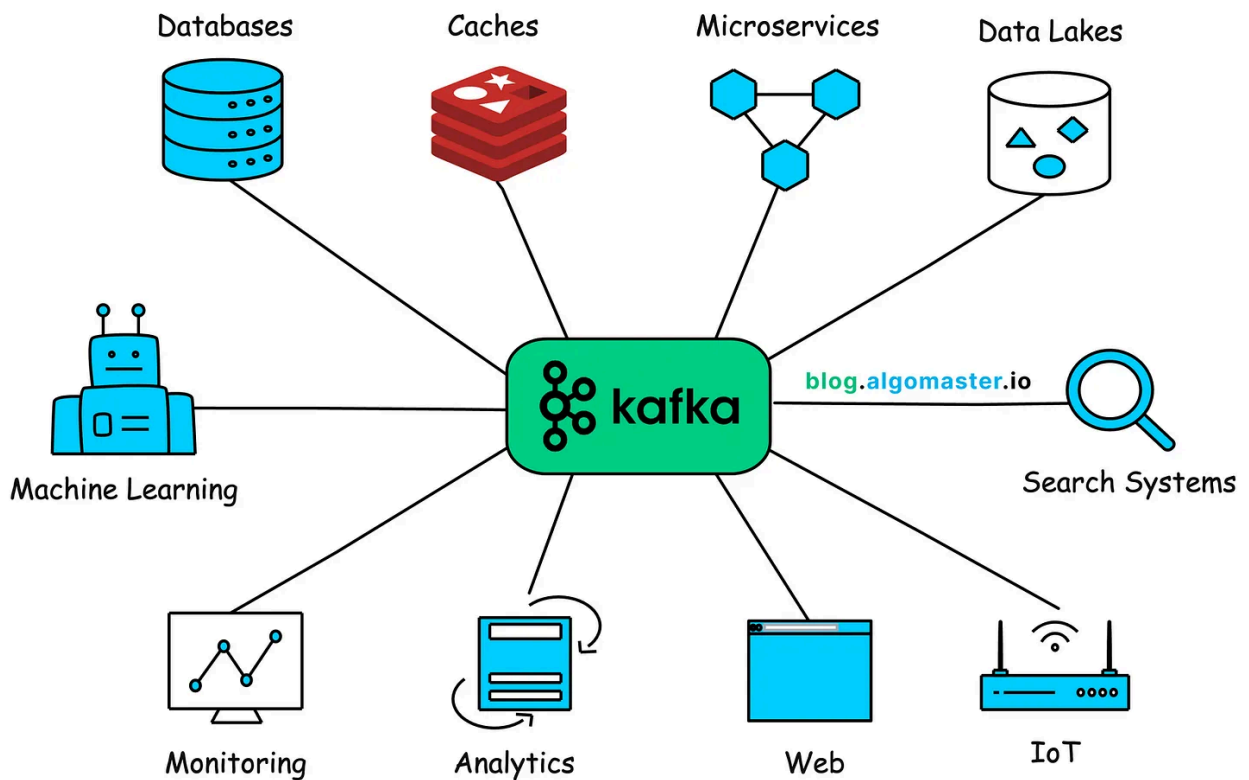
1

9

Share

Apache Kafka began its journey at **LinkedIn** as an internal tool designed to collect and process massive amounts of log data efficiently. But over the years, Kafka has evolved far beyond that initial use case.

Today, Kafka is a **powerful, distributed event streaming platform** used by companies across every industry—from tech giants like Netflix and Uber to banks, retailers, and IoT platforms.



Its core architecture, based on **immutable append-only logs**, **partitioned topics**, and **configurable retention**, makes it incredibly scalable, fault-tolerant, and versatile.

In this article, we'll explore the **10 powerful use cases of Kafka** with real-world examples.

1. Log Aggregation

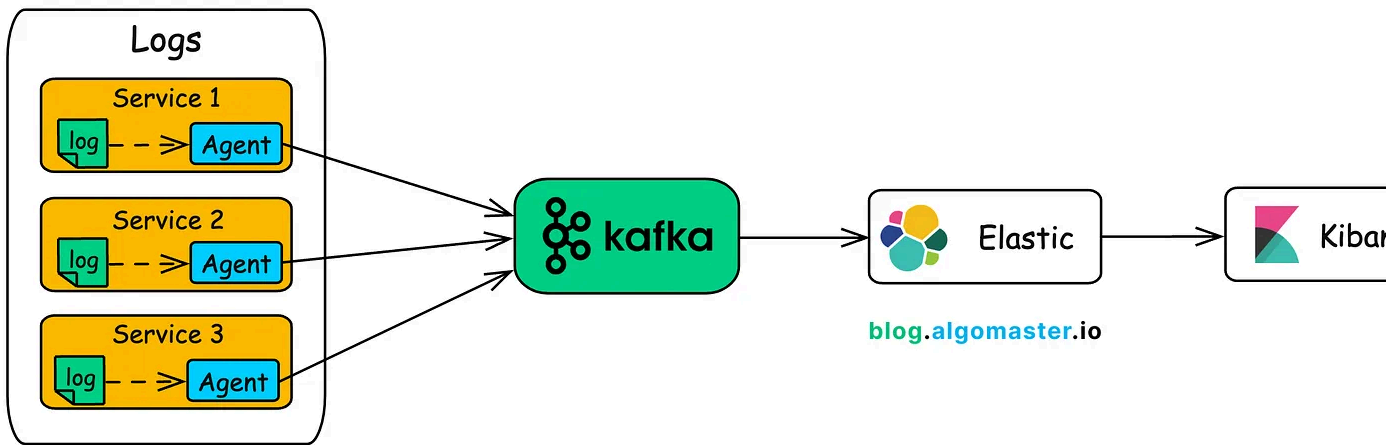
In modern distributed applications, logs and metrics are generated across hundreds of servers, containers, and applications. These logs need to be collected for monitoring, debugging, and security auditing.

Traditionally, logs were stored locally on servers, making it difficult to search, correlate, and analyze system-wide events.

Kafka solves this by acting as a **centralized, real-time log aggregator**, enabling fault tolerant, scalable, and high-throughput pipeline for log collection processing.

Instead of sending logs directly to a storage system, applications and logging agent stream log events to Kafka topics. Kafka provides a durable buffer that absorbs spikes in log volumes while decoupling producers and consumers.

Kafka Log Aggregation Pipeline



Step 1: Applications Send Logs to Kafka (Producers)

Each microservice, web server, or application container generates logs in real time sends them to Kafka via lightweight log forwarders (log agents) like: **Fluentd**, **Logst** or **Filebeat**.

These tools publish logs to specific Kafka topics (e.g., `app_logs`, `error_logs`).

Step 2: Kafka Brokers Store Logs

Kafka acts as the central, durable and distributed log store, providing:

- **High availability** - logs are replicated across multiple brokers
- **Persistence** - logs are stored on disk for configurable retention periods
- **Scalability** - Kafka can handle logs from thousands of sources

Step 3: Consumers Process Logs

Log consumers (like Elasticsearch, Hadoop, or cloud storage systems) read data from Kafka and process it for:

- **Indexing** - for searching and filtering logs
- **Storage** - long-term archiving in S3, HDFS, or object storage
- **Real-time monitoring** - trigger alerts based on log patterns

Step 4: Visualization and Alerting

Processed logs are visualized and monitored using tools like:

- **Kibana / Grafana** - for dashboards and visualization
- **Prometheus / Datadog** - for real-time alerting
- **Splunk** - for advanced log analysis and security insights

2. Change Data Capture (CDC)

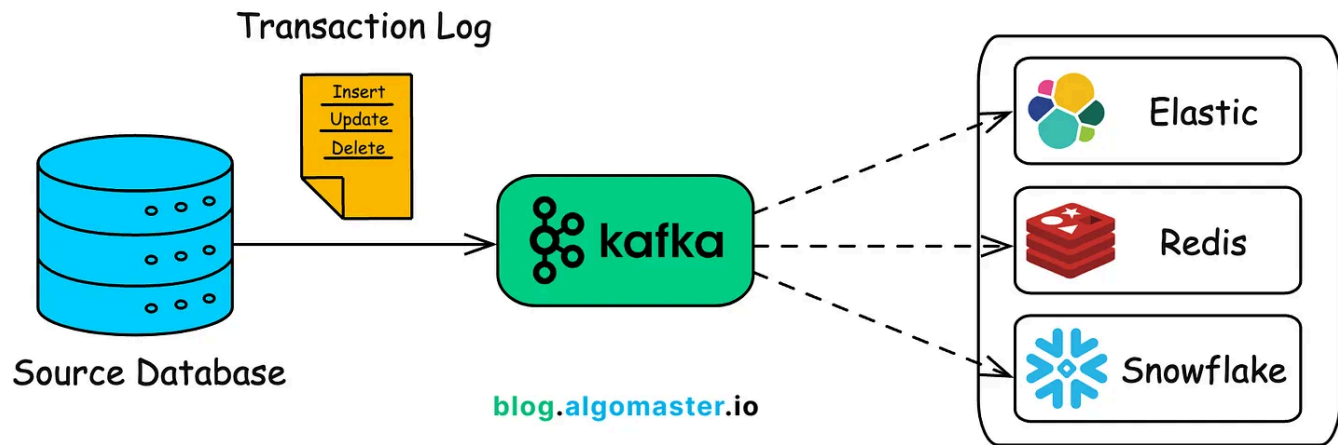
Change Data Capture (CDC) is a technique used to **track changes in a database (inserts, updates, deletes)** and **stream those changes in real time** to downstream systems.

Modern architectures rely on multiple systems—search engines, caches, data lakes, microservices—all of which need **up-to-date data**. Traditional **batch ETL jobs** are slow, introduce latency, and often lead to:

- **Stale data** in search indexes and analytics dashboards
- **Inconsistencies** when syncing multiple systems
- **Performance overhead** due to frequent polling

Kafka provides a **high-throughput, real-time event pipeline** that captures and distributes changes from a source database to multiple consumers—ensuring low latency and consistency across systems.

How CDC Works with Kafka



Step 1: Capture Changes from the Database

Tools like Debezium, Maxwell, or Kafka Connect read the database transaction log (binlogs, WALs) to detect: INSERTs, UPDATEs, DELETEs

Each change is transformed into a structured event and published to a Kafka topic.

Step 2: Stream Events via Kafka

Kafka topics act as an **immutable commit log**, providing:

- **Durability** — All changes are stored reliably
- **Ordering** — Events for a given key (e.g., primary key) are strictly ordered
- **Scalability** — Thousands of events per second per partition

Step 3: Distribute to Real-Time Consumers

Multiple consumers can subscribe to the change events for various use cases:

- **Search Indexing** → Sync changes to Elasticsearch / OpenSearch
- **Caching** → Update Redis / Memcached for fast reads
- **Analytics** → Stream into BigQuery / Snowflake / Redshift

3. Event-Driven Microservice Communication

In a **microservices architecture**, dozens or hundreds of services must coordinate to fulfill business processes.

Traditional **REST-based communication** introduces **tight coupling**, meaning:

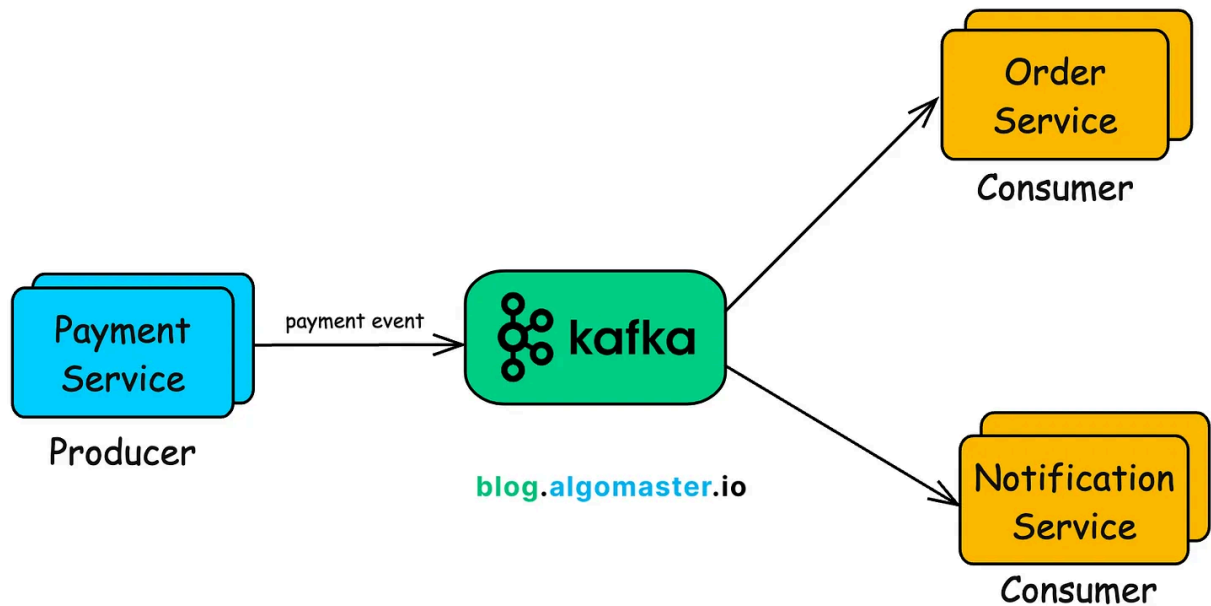
- Services depend directly on each other, creating **cascading failures** if one goes down.
- Scaling becomes **challenging**, as each service call adds latency and load.
- If a service is **temporarily unavailable**, messages are **lost**, requiring retries.

Kafka enables **event-driven microservices**, where:

1. **Services publish events** to Kafka topics (e.g., `payment_completed`) instead of calling each other directly.
2. **Other services subscribe** to relevant events and react when needed.
3. **Kafka buffers messages**, ensuring no data loss if a service is temporarily unavailable.
4. **Services scale independently**, with Kafka **load-balancing messages** across instances.

This creates a loosely-coupled, resilient architecture where services **communicate asynchronously via events**.

Example Workflow: Payment Event



When a customer makes a payment:

- **Payment Service** → publishes a `PaymentCompleted` event to Kafka
- **Order Service** → consumes the event and updates the order status
- **Notification Service** → listens to the event and sends an email receipt

Each service operates **independently**, reacting to events. If a service is down, Kafka retains the event until it is available.

Many companies, including Netflix, Uber, and Airbnb, use Kafka to build highly scalable, resilient microservices.

4. Real-Time Machine Learning Pipeline

Today's businesses generate massive volumes of data continuously—from user interactions and system logs to sensor readings and financial transactions. To remain competitive, companies must extract insights and make decisions as fast as the data arrives.

Traditional ML pipelines rely on **batch processing**—ingesting data periodically, training models offline, and applying predictions in delayed cycles. While this works for some use cases, it fails in scenarios that require **immediate feedback, continuous adaptation, and low-latency predictions**.

Limitations of Traditional ML Pipelines

- **Delayed predictions** — Batch jobs run every few hours or days
- **Stale insights** — Models act on outdated data
- **Slow feedback loops** — Difficult to adapt to fast-changing patterns
- **Limited responsiveness** — Inadequate for time-sensitive decisions like fraud detection or real-time recommendations

Kafka addresses these challenges by acting as a **real-time data pipeline**—moving data between producers, processors, ML models, and consumers with **low latency, high throughput, and fault tolerance**.

Kafka enables ML systems to:

- **Ingest live data streams** from apps, sensors, and services
- **Enrich, clean, and transform data** in real time
- **Trigger model inference** instantly
- **Continuously update models** with fresh data
- **Feed predictions** to downstream systems for immediate action

How Kafka Powers Real-Time ML Pipelines

Step 1: Live Data Ingestion

Kafka collects real-time data from various sources:

- **User interactions** (clicks, scrolls, purchases)

- **System metrics** (errors, latency, usage patterns)
- **IoT devices** (temperature, location, movement)
- **Transactions** (payments, transfers)

Each event is published to a Kafka topic (e.g., `user_events`, `transactions`, `sensor_data`), creating a central stream of real-time inputs.

Step 2: Feature Extraction & Preprocessing

Stream processors like **Kafka Streams**, **Apache Flink**, or **Spark Structured Streaming** consume Kafka topics and:

- Clean and normalize incoming data
- Perform **feature engineering** (e.g., encode behavior into numerical vectors)
- Aggregate data within time windows
- Join with reference data (e.g., user profiles, product catalogs)

The resulting enriched features are either:

- Sent to **feature stores** (e.g., Feast)
- Passed directly to **real-time model inference** services

Step 3: Real-Time Inference

Trained machine learning models make predictions in real time by consuming feature streams:

- **Fraud detection:** Flag a suspicious payment instantly
- **Product recommendations:** Suggest items while the user browses
- **Churn prediction:** Detect user dissatisfaction before they leave
- **Autonomous vehicles:** Make navigation decisions in milliseconds

Predictions are published to new Kafka topics (e.g., `fraud_scores`, `recommended_items`) for further processing or immediate action.

Step 4: Feedback & Model Retraining

Kafka also enables **continuous learning**:

- Feedback data (e.g., user reactions, outcomes, labels) is streamed to Kafka
- Training pipelines consume this data to **retrain or fine-tune models**
- Periodically retrained models are **redeployed** into the inference layer

This tight feedback loop allows systems to stay relevant and adapt to new patterns quickly.

Step 5: Serving Predictions & Taking Action

Predictions are routed to different consumers via Kafka:

- **Dashboards** for human insights
- **Microservices** to trigger business logic (e.g., show an offer, block a transaction)
- **Databases or data lakes** for long-term storage and analysis
- **Alerting systems** for automated monitoring and decisioning

5. Real-Time Clickstream Analysis

Understanding how users interact with websites and mobile apps is **mission-critical** for e-commerce, content, and advertising platforms. Every **click, scroll, search**, or **view** offers valuable behavioral insights.

Traditional Approach: Too Little, Too Late

- Clickstream data was often logged and processed **hours or days later**

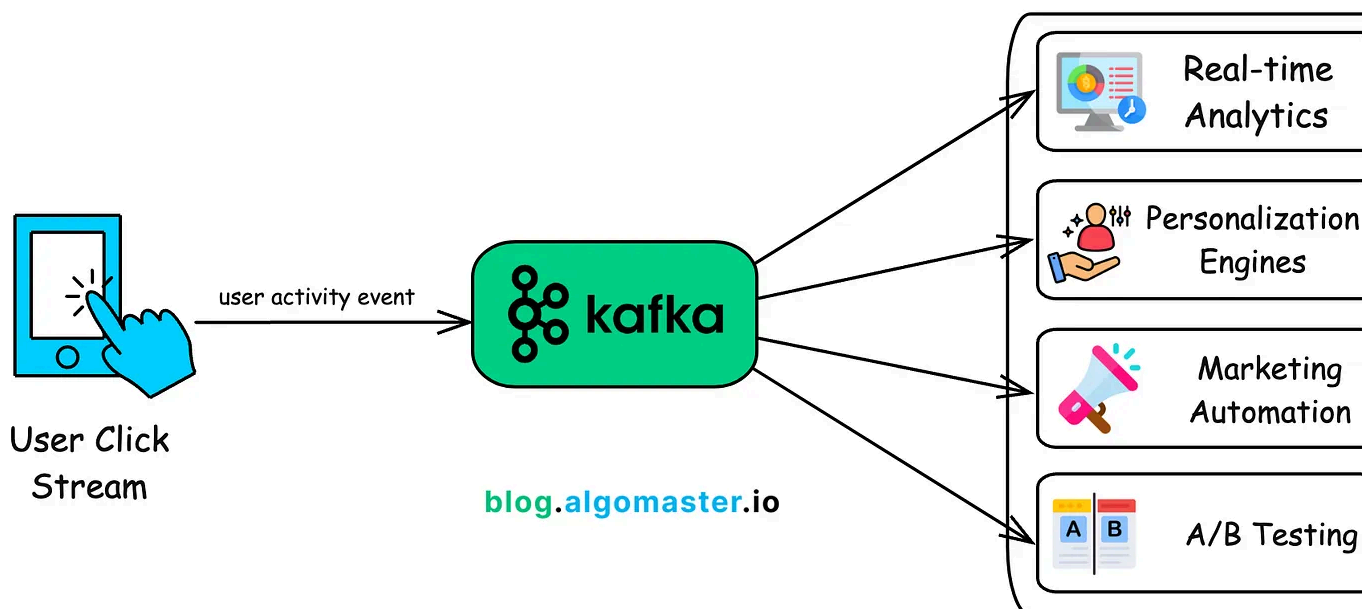
- This delayed **personalization, retargeting, and UI optimizations**
- High user volumes made it difficult to process interactions at scale

As a result, opportunities to **react while the user is still engaged** were missed.

Kafka provides the backbone for modern clickstream analytics, enabling:

- **Low-latency ingestion** of millions of events per second
- **Durable storage and ordering** for user session data
- **Parallel processing** by multiple downstream systems
- **Real-time personalization and decision-making**

How it Works



1. User Interaction → Kafka Topic

Every time a user:

- Clicks a product
- Scrolls a page

- **Likes** a post
- **Searches** for content

...the frontend app captures the event and publishes it to a Kafka topic (e.g., `user_activity`).

Each event includes:

- `user_id`
- `timestamp`
- `action_type` (click, scroll, search, etc.)
- `item_id` (product, article, ad)
- Metadata (device type, location, referrer, etc.)

2. Kafka Buffers and Distributes Events

Kafka:

- **Persists all events in order** (partitioned by user/session ID)
- Supports both **real-time streaming** and **batch ETL**
- Acts as a **durable buffer**, so slow consumers don't lose data

3. Multiple Consumers Process Clickstream Data

- **Real-Time Analytics:** Track traffic spikes, top pages, bounce rates on live dashboards
- **Personalization Engines:** Recommend products or content based on recent clicks/searches
- **Marketing Automation:** Trigger in-session offers, push notifications, or retargeting ads
- **A/B Testing:** Compare performance of design/layout variants in real time

- **AdTech Systems:** Measure ad impressions, clicks, and attribution instantly

6. IoT Data Ingestion

In IoT (Internet of Things) applications, massive networks of sensors and devices generate continuous streams of data—from **machines on factory floors**, to **smart home appliances**, to **connected vehicles**.

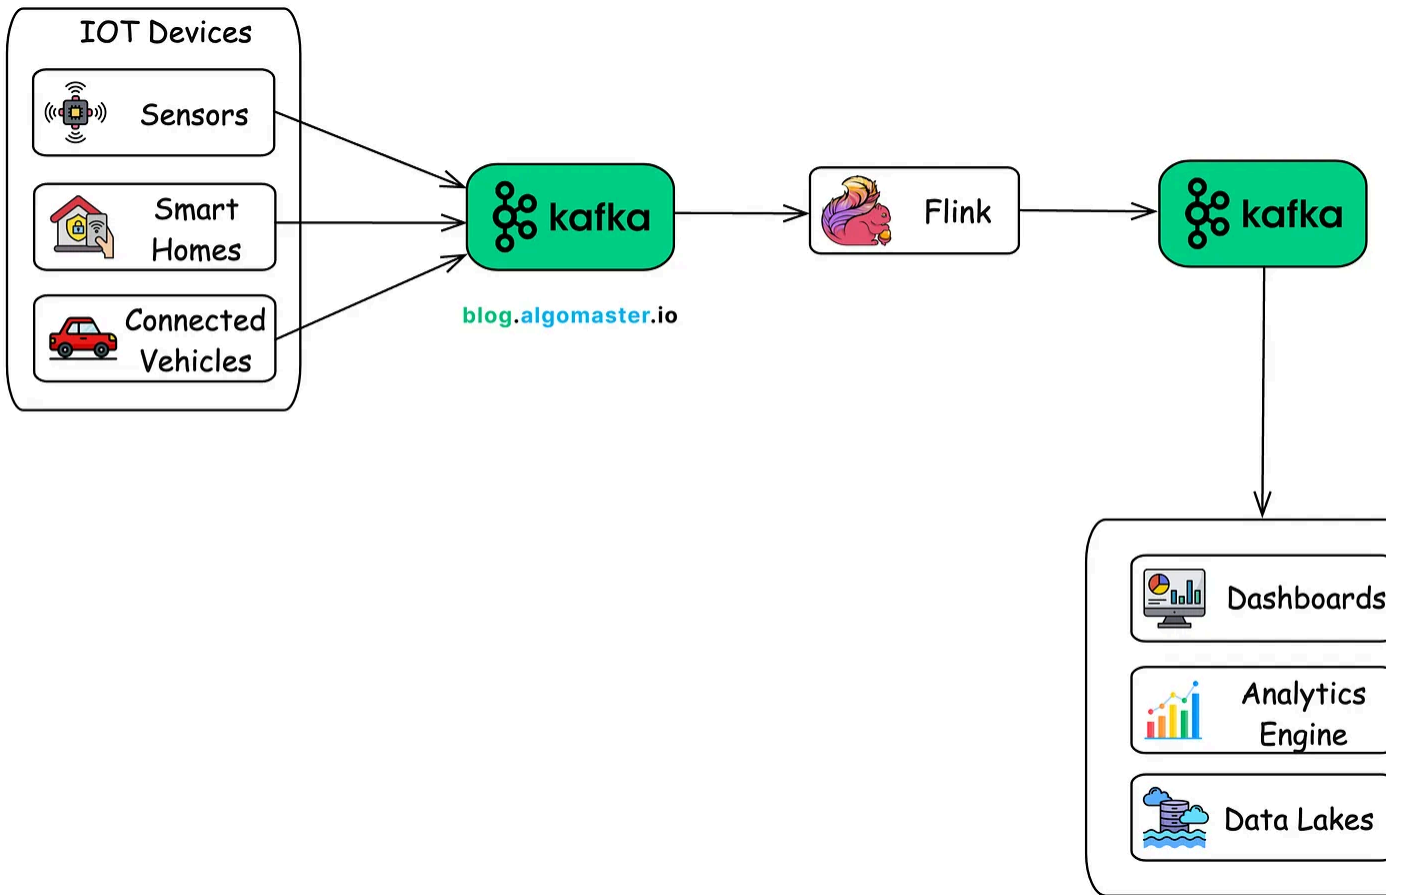
Collecting, processing, and acting on this data in real time is essential, but challenging due to:

- **High event volume** (millions of devices sending readings every second)
- **Unreliable network connections**
- **The need for ordered, time-series data**
- **Requirements for low-latency processing and automated control**

Apache Kafka serves as the **central nervous system** for ingesting and distributing IoT data. It offers:

- **High-throughput ingestion** from millions of devices
- **Fault-tolerant buffering** in case of network interruptions
- **Message ordering** within each partition (keyed by device ID)
- **Scalability** by simply adding more brokers
- **Integration with IoT protocols** like MQTT via Kafka Connect

How Kafka Powers IoT Data Pipelines



Step 1: Devices Produce Data

Each IoT device acts as a **Kafka producer**—either directly or via an IoT gateway that batches and forwards data to Kafka.

- Example: A smart agriculture deployment streams **temperature, humidity, and soil moisture** readings from distributed sensors.
- Kafka topics can be organized by **sensor type, location, or device ID**.

Step 2: Kafka Buffers & Orders Data

Kafka brokers ingest data and retain it durably, even if:

- A device temporarily goes offline
- Network connectivity is unstable
- The consumer system is lagging behind

Kafka ensures event ordering per device using partitions, which is critical for time series analysis and real-time control systems.

Step 3: Real-Time Processing & Control

Kafka Streams or external stream processors (e.g., Apache Flink, Spark Streaming) c

- Detect anomalies (e.g., vibration spikes on industrial machines)
- Aggregate sensor data to assess device or system health
- Trigger actions (e.g., send maintenance alerts or shut down machines)

Step 4: Fan-Out to Consumers

Kafka distributes IoT data to:

- Dashboards for live device status
- Analytics engines for trend and anomaly detection
- Data lakes for model training and historical analysis
- Control systems that send commands back to actuators via Kafka topics

7. Payment Processing Pipelines

Modern payment systems—whether it's online checkouts, mobile wallets, or instant bank transfers—demand fast, reliable, and ordered event processing across multiple services.

From the moment a user clicks “Pay” to the final confirmation and receipt, multiple systems must work together in real time:

- Payment gateways
- Credit/debit processors

- Fraud detection services
- Order and inventory management
- Accounting and ledger systems
- Notification and customer service systems

Payment pipelines must:

- **Process events in strict order** (e.g., payment before confirmation)
- **Ensure delivery without duplication or data loss**
- **Handle high spikes in volume** (e.g., flash sales, festival traffic)
- **Recover gracefully** from service failures without losing transactions
- **Maintain consistency** across multiple services

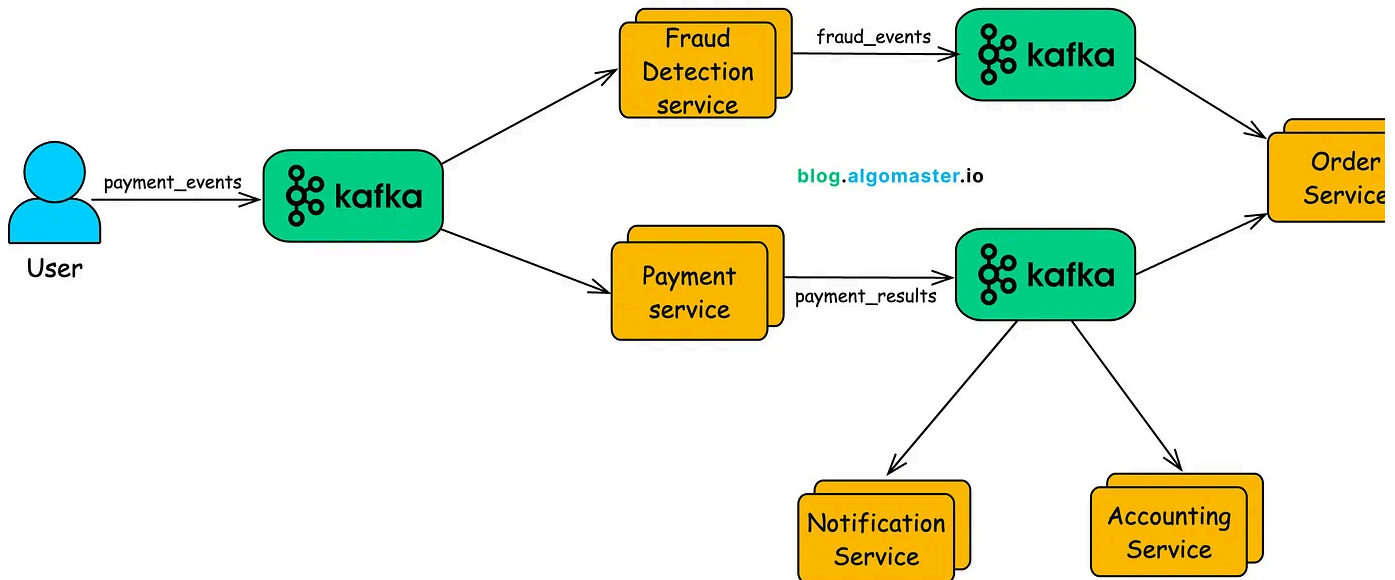
Traditional point-to-point REST architectures struggle with these requirements. Failures or delays in one service can ripple across the system.

Kafka provides the ideal backbone for real-time payment pipelines by offering:

- **Reliable event delivery**
- **Message ordering per transaction or order**
- **At-least-once or exactly-once semantics**
- **High throughput to handle burst loads**
- **Decoupled, resilient architecture**

Instead of direct service-to-service calls, **payment workflows are modeled as a chain of events** published to and consumed from Kafka topics.

Example Workflow: E-Commerce Payment



1. **User initiates payment:** PaymentInitiated event published to Kafka (keyed by Order ID)
2. **Payment service:** Consumes event, attempts payment, publishes PaymentSuccessful or PaymentFailed
3. **Fraud detection service:** Subscribes to PaymentInitiated, runs checks, publishes FraudCheckPassed or FraudDetected
4. **Order service:** Subscribes to PaymentSuccessful and FraudCheckPassed, updates order status
5. **Notification service:** Sends confirmation email or failure alert based on event outcome
6. **Accounting and audit systems:** Stream payment results into databases or data lakes for compliance and reporting

Kafka's idempotent producers and transactional consumers help prevent:

- Duplicate payment attempts
- Out-of-order state transitions
- Data corruption in ledgers or audit logs

Using Kafka's **exactly-once semantics**, systems can safely process financial transactions even in complex, multi-step workflows.

8. Stateful Stream Processing

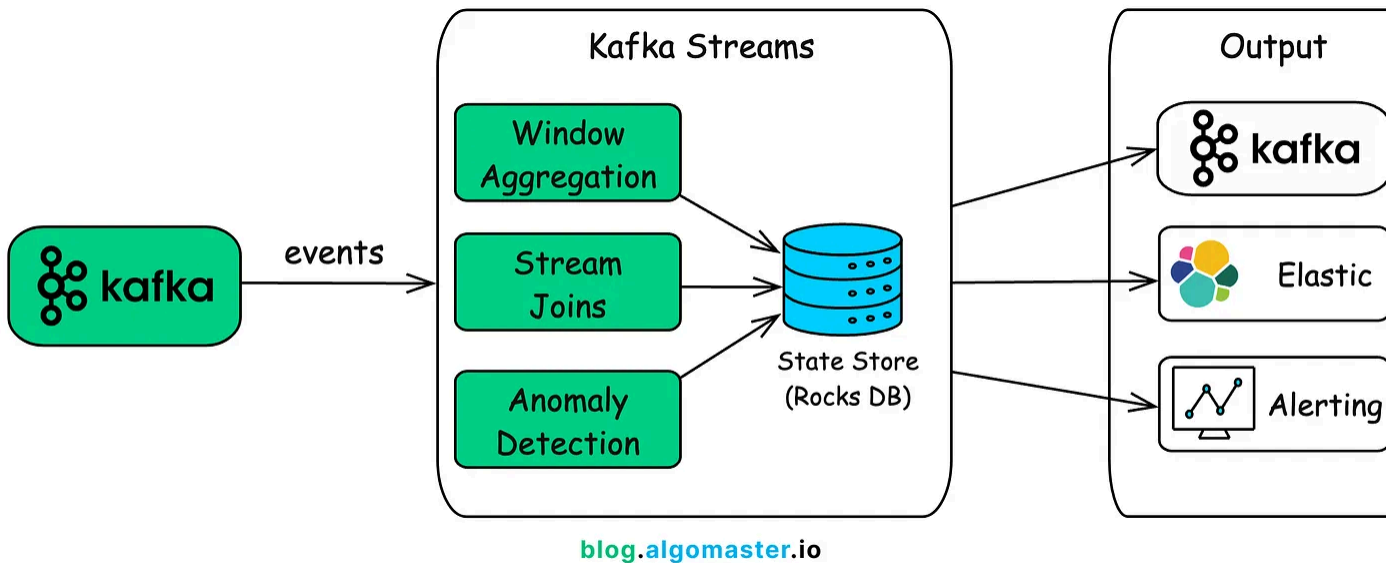
Traditional stream processing focuses on handling **independent events**, but many real-world use cases require **stateful operations**—where **past events influence current computations**.

Key challenges include:

- **Maintaining context** across events (e.g., tracking a user's session).
- **Handling windowed aggregations** (e.g., computing moving averages).
- **Processing event sequences correctly** (e.g., detecting fraud based on past transactions).

Kafka Streams enables **stateful event processing** by maintaining **local state stores** to track event history without needing an external database.

How It Works:



1. Event Ingestion & Partitioning

Kafka streams data from topics partitioned by key (e.g., `user_id`, `session_id`, `account_id`).

- Events from the same key always go to the **same stream processor instance**
- This ensures all relevant history is **co-located with the processor**

2. Maintaining Local State

Kafka Streams uses **embedded state stores** (backed by RocksDB) to track:

- User sessions
- Running aggregates
- Recent event patterns
- Time-windowed activity

These state stores are **local to each stream task** and **automatically backed up** to Kafka for recovery.

3. Real-Time Stateful Computations

Examples of stateful operations include:

- **Windowed Aggregations:**
 - Count page views per session
 - Compute moving averages over the last 10 minutes
- **Stream Joins:**
 - Enrich clickstream events with user profile data
 - Combine payment events with fraud scores
- **Pattern & Anomaly Detection:**
 - Detect login attempts from multiple locations
 - Identify outlier transactions based on past behavior

4. Output & Action

After processing, results can be:

- Published back to Kafka topics
- Written to external systems (e.g., Elasticsearch, PostgreSQL, Redis)
- Sent to monitoring/alerting tools (e.g., trigger a fraud alert or anomaly dashboard)

9. Real-Time Fraud and Anomaly Detection

In the financial world, **speed is everything** when it comes to detecting fraud. Whether it's credit card abuse, identity theft, or suspicious withdrawals, delays in identifying malicious activity can lead to:

- **Massive financial losses**

- **Erosion of customer trust**
- **Regulatory penalties**

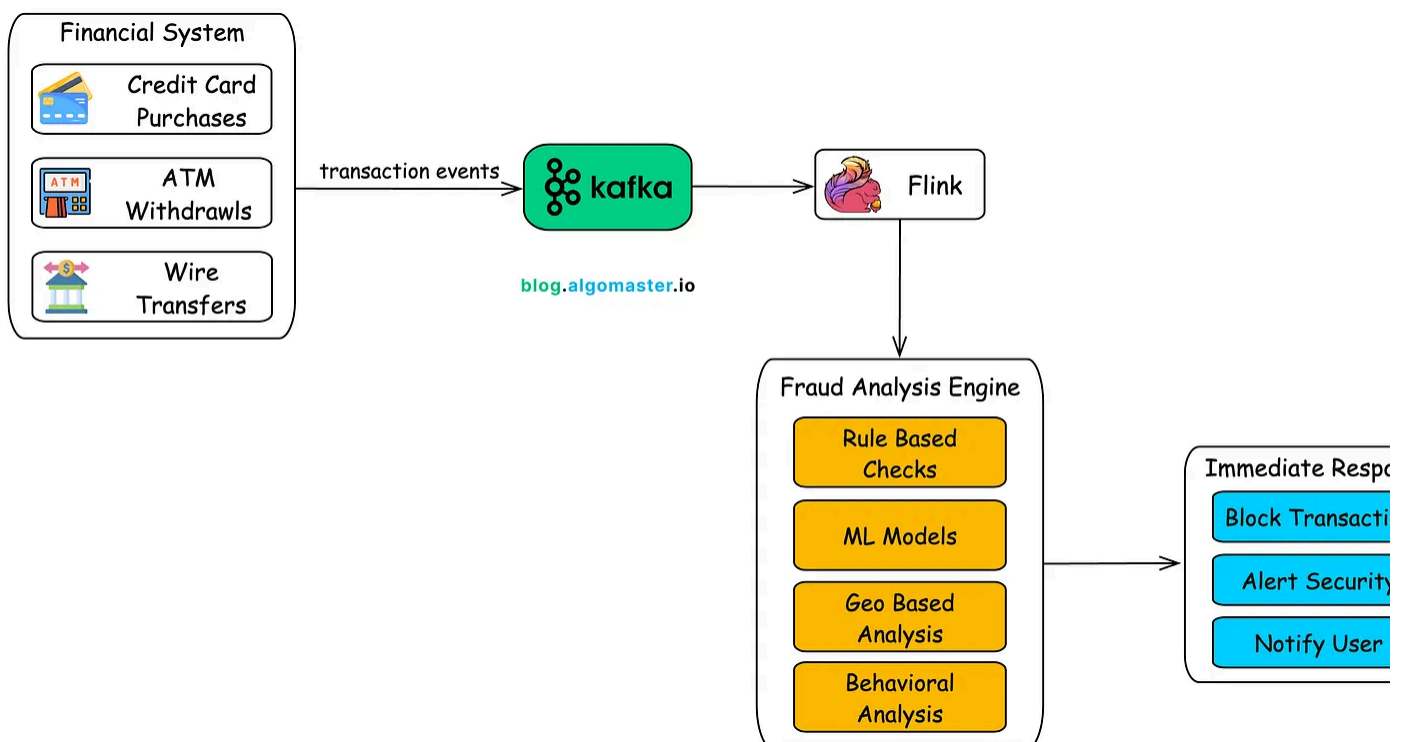
Traditional fraud detection systems, which rely on **batch processing**, often identify threats **after the damage is done**.

Kafka changes this by enabling **real-time detection and prevention**.

Kafka acts as the backbone of a **streaming fraud detection pipeline**, providing:

- **Low-latency ingestion** of transaction data
- **Scalable parallel processing** using partitions (e.g., by account ID)
- **Durability and ordering** to ensure consistency and auditability
- **Integration with ML/stream processing frameworks** like Kafka Streams, Flink and Spark Streaming

How It Works:



1. Transaction Event Ingestion

- Financial systems stream **transaction events** into Kafka topics (e.g., `transaction_events`)
- Events are **partitioned by key** (e.g., account ID, user ID)
- Kafka ensures all events for a given account are processed **in order**, enabling time-based analysis

2. Real-Time Fraud Analysis

Fraud detection logic runs as stream processors (Kafka Streams, Apache Flink, etc.) that analyze live data:

- **Rule-based checks** (e.g., withdrawal limits, velocity checks)
- **ML models** trained on historical data
- **Geo-based filters** (e.g., impossible travel between ATM transactions)
- **Behavioral anomaly detection**

These systems compare **real-time data with historical patterns** to flag suspicious behavior immediately.

3. Immediate Response

When a transaction is flagged as suspicious:

- It can be **blocked or held** for manual review
- **Alerts** are triggered for security teams
- The user may receive a **real-time notification** (e.g., SMS, email, or in-app prompt)

Kafka ensures that even during **high-traffic periods** (like holiday seasons), no transaction slips through unnoticed.

10. Media Streaming Analytics

Modern media platforms—such as **Netflix**, **Spotify** and **YouTube**—serve millions of users daily across the globe. These platforms rely on real-time data to:

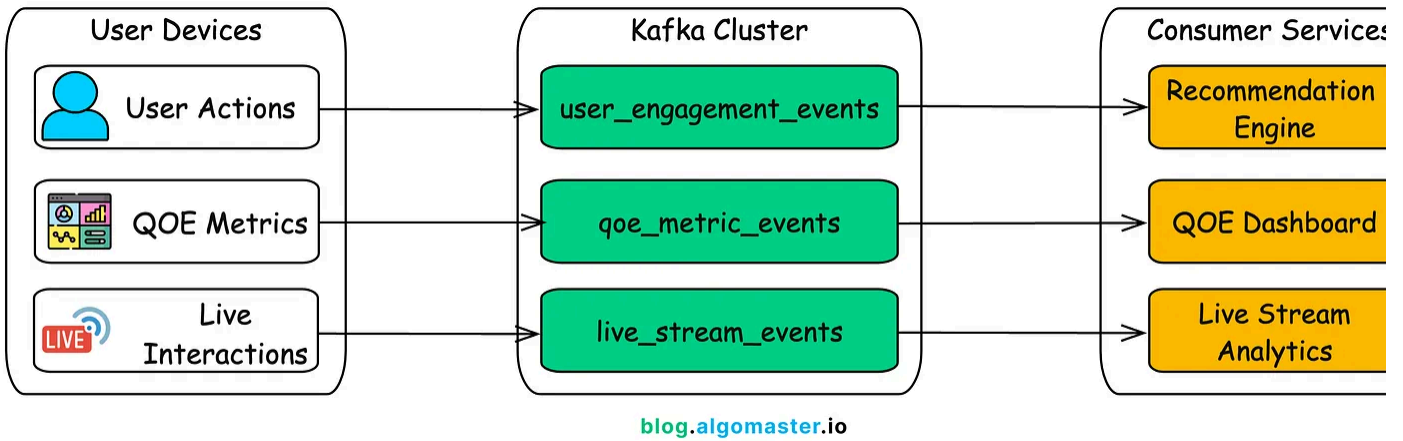
- Personalize content recommendations
- Monitor video/audio playback quality
- Track engagement with live events
- Measure content performance instantly

Handling this constant firehose of **user interactions**, **system metrics**, and **content insights** is a big data problem—in real time.

Apache Kafka is the **central data highway** powering real-time analytics in the media industry. It enables platforms to:

- Collect **millions of events per second**
- Maintain **event ordering** (e.g., per user or session)
- **Buffer and route data** to downstream analytics and ML systems
- Decouple **user-facing applications** from **backend analytics** to ensure smooth performance

Media Streaming Use Cases Powered by Kafka



1. User Engagement Tracking

Every user action—play, pause, rewind, skip, rate—is published as an event to Kafka

- UserX watched VideoY for 7 minutes
- UserY paused VideoZ at 3:42

These events feed into:

- **Recommendation engines** → to personalize the next piece of content
- **Trending content services** → e.g., “Top 10 Today” lists
- **A/B testing pipelines** → to compare UI layouts or algorithm variants in real time

2. Quality of Experience Monitoring

Video/audio players send periodic **QoE metrics** (e.g., buffering events, bitrate switches, playback errors) to Kafka topics.

Kafka enables:

- Real-time dashboards for operations teams
- Alerting systems to detect regional or CDN-specific issues
- Automated remediation logic (e.g., switching CDNs)

3. Live Streaming Interactions

During live events (concerts, sports, streams), Kafka powers real-time features:

- **Live chat and emoji reactions**
- **Viewer counts**
- **Sentiment analysis** for engagement monitoring

Kafka distributes these events to all interested consumers—analytics systems, moderation services, dashboards—**instantly and at scale**.

4. Ad Insertion

For ad-supported content, Kafka streams user context (device, location, behavior) to an **ad decision service**, which returns the best ad to play next—**all within milliseconds**.

Kafka ensures:

- **Low-latency decisioning**
- **No data loss during peak loads**
- **Fine-grained targeting and real-time performance tracking**

5. Content Performance and Feedback Loop

Kafka feeds aggregated data to:

- **Business intelligence dashboards** (e.g., top-performing titles, watch duration stats)
- **Editorial teams** → to promote trending content
- **ML models** → for improving personalization and retention

Thank you for reading!

If you found it valuable, hit a like ❤️ and if you have any questions or suggestions, leave a comment.

I hope you have a lovely day!

See you soon,

Ashish



78 Likes · 9 Restacks

← Previous

Next →

Discussion about this post

Comments

Restacks



Write a comment...



Srikanth shoda Jun 26

❤️ Liked by Ashish Pratap Singh

Thanks, for sharing such an insight blog on different use cases of Kafka!

❤️ LIKE (1) 💬 REPLY



© 2025 Ashish Pratap Singh · [Privacy](#) · [Terms](#) · [Collection notice](#)
[Substack](#) is the home for great culture